

Lab 3 - CUDA

NGUYEN Duc Tung

I implement the CUDA version with following steps:

- Prepare device's memory (allocation, memory copy)
- Calculating block size
- Writing the kernel and invoke it
- Copy output back from device to host
- Free device's memory

Here is the speed up:

- CPU version: 2800 ms
- CPU - OpenMP version: 700 ms
- CUDA version: 21.8 ms